

# Decentralization

- Foundation Crypto
  - Public Key Systems
    - RSA
    - ECDSA
  - Hashing
    - Characteristics
    - Algorithms (MD5, SHA-X)
    - Use as PoW
- P2P systems – Napster vs Gnutella
- Blockchain – is this before or after IPFS
  - Bitcoin intro – transaction, blocks, mining, distributed consensus
  - Ethereum – smart contracts
  - ERC-20 Tokens
- IPFS
  - Self-Certified Identities
  - Content-Addressable block store
- ERC-20 Tokens
- Filecoin

# Decentralized Computing

- Involves a system where the hardware and software resources are generally *Distributed* across a network
- Location of Control is a key part of Decentralized Computing
  - There should be no central point of control
  - Often you apply a test: If a very powerful entity wanted to shut down the system – could they do it? How easily?
    - Impossible : Highly decentralized
    - Difficult: partially decentralized
- Often involves groups of users contributing resources to be shared by all
  - Content
  - Software
  - Compute capability
  - Storage

# Napster

- Classic example in Computing:
- released in 1999 by Shawn Fanning/Sean Parker
- Enabled global (illegal) distribution of copyrighted material (music/video)
- Operated over P2P – so partially decentralized
- Needed Napster.com to initially link consumer with suppliers
- Court Order shutdown the service in 2001
- Followed by systems like Gnutella
- Limewire was a Gnutella client – came under court scrutiny – tried to make their client legal – inserted code into version 5.5.10 and later that allowed them to cut users off from the network -

# Decentralized Systems: Pros/Cons/Ethics

- Some Pros

- Makes for highly resilient systems than can survive outages
- Often makes very efficient use of resources
- Denies special position for the controlling party
- Censorship Resistance : (Good for freedom fighters etc)

- Some Cons

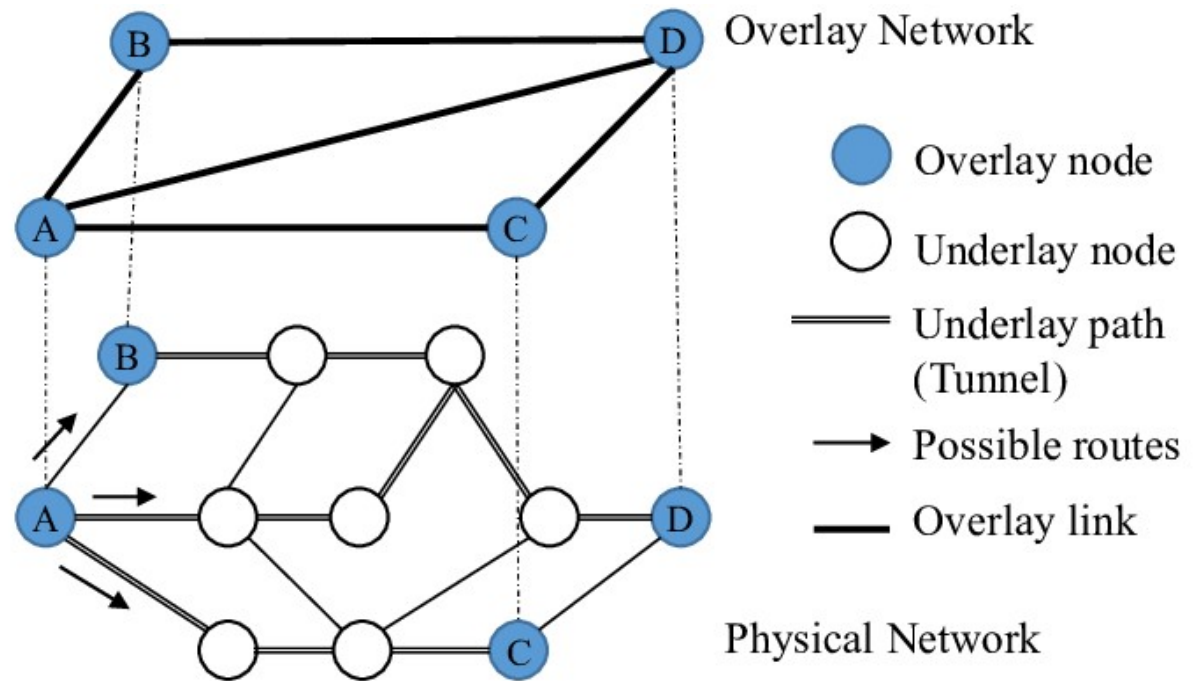
- Sometimes very inefficient or resource intensive (at one point P2P traffic with illegal content occupied > 50% of bandwidth on Internet Trunks)
- Subject to the 'Tragedy of the Commons' – no incentives to manage common resource pool wisely
- Difficult to Incentivize people to work towards the common good
- Censorship Resistance : (Good for Child Pornographers, Drug dealers, Revenge Porn)
  - Individuals may not like Government control but societies usually dox

# Key Technologies

- Peer-to-Peer (P2P) networking to provide transport & discovery
- Some key crypto technologies
  - Hashing
  - Distributed Hash Tables (DHTs)
  - Public/Private Keys
  - Signatures
  - Merkle Trees
- Higher Level abstractions built on these
  - Blockchains
  - IPFS


# Peer to Peer Networking - Gnutella

- A node wishing to join the Gnutella Network (somehow) finds some node that is already on it
  - Some peers may be hard-coded into the client; nodes can refuse connection and re-direct
- Leads to the establishment of an Overlay Network – a set of nodes communicating over a mesh of single point-to-point IP connections



# Gnutella Primitives


- Used as a distributed search protocol
- Messages
  - Ping – used to discover hosts
  - Pong – response to ping with client address
- Network construction showed on next slide
- Very Resilient
  - Node failure causes neighbours to seek new peers
  - Difficult for authorities to shut down
  - Difficult for network operators to try to manage – even when they are trying to be helpful to it!



# Joining Gnutella Network

- The new node connects to a well known 'Anchor' node.
- Then sends a PING message to discover other nodes.
- PONG messages are sent in reply from hosts offering new connections with the new node.
- Direct connections are then made to the newly discovered nodes.

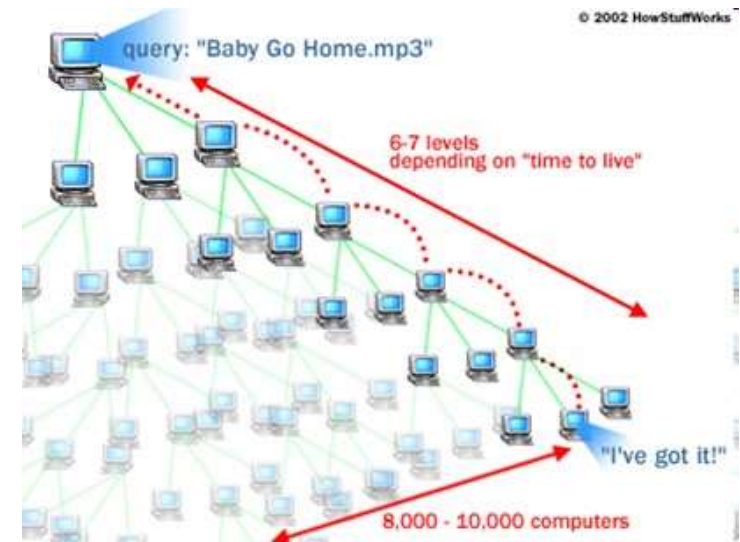
## Gnutella Network





# Searching the P2P network

- Query – search mechanism
- QueryHit – response to query, include info necessary to get data
- Finding content
  - Ping/Query sent to all connected clients
  - Pong/QueryHit sent back along return path
  - TTL mechanism to limit distance
  - File downloads by Http on direct connection



# Uses of P2P Networks

- Can be constructed for many purposes
  - Content distribution (Bittorrent)
  - Distribution of transactions (Bitcoin, Ethereum)
  - Storage of information (DHT's mentioned later)
  - File Sharing (InterPlanetary File System(IPFS), Swarm)
  - Person-to-Person Messaging (Whisper)

# Fundamental Technologies – Cryptographic Hash Functions

- Checksums have been used for years to detect errors in blobs of data
- Iterate over the bytes in a blob => fixed size checksum
- Hash functions were invented (sometime in the 1950/60s) along with the concept of a hash table. A hash function
  - Operates on the “Key” field of a hash tables to produce a fixed size result
  - People invent their own hash functions – it is desirable that the hash function ‘spreads’ well over the range of possible outputs

# Cryptographic Hash Functions

Fox	cryptographic hash function	DFCD 3454 BBEA 788A 751A 696C 2409 7009 CA99 2D17
The red fox jumps over the blue dog	cryptographic hash function	0086 46BB FB7D CB22 823C ACC7 6CD1 90B1 EE6E 3ABC
The red fox jumps over the blue dog	cryptographic hash function	8FB8 7558 7851 4F32 D1C6 76B1 79A9 0DA4 AEEF 4819
The red fox jumps over the blue dog	cryptographic hash function	FCD3 77DB 5AF2 C6FF 315F D401 C0A9 7D9A 46AF FB45
The red fox jumps over the blue dog	cryptographic hash function	8ACA D682 D588 4C75 4BF4 1799 7D88 BCF8 32B9 6A6C

- Iterate over a blob of data to produce a fixed size result (message digest)
  - Deterministic – always produces the same output for the same message
  - Quick to compute
  - Infeasible to generate a message that yields a given hash value
  - Infeasible to find two messages with the same hash value
  - A small change in the message should change the hash value dramatically (avalanche effect)
- Note that a hash of a document (file, block, movie) etc uniquely identifies that document – can be used in a document database to retrieve content.
- Used to prove document integrity – If you can compare to a trusted hash – you know the document has not been altered
- Can be used to commit to things that have not yet been revealed (later revealing a 'pre-image' of the hash)


# Evolution of Hash Functions

- Cryptographic primitives get attacked! – If any flaws are found, replacements are designed
  - Message Digest 5 (MD5) designed by Ron Rivest in 1991 – produces a 128-bit digest – considered broken
  - SHA-1 – US government NIST in 1995 - 160 bit digest – considered broken
  - Sha-2- consists of Sha-256 & SHA-512 – worries about SHA-256, but in widespread usage
  - SHA-3 – US Government NIST in 2015 – based on Keccak primitives



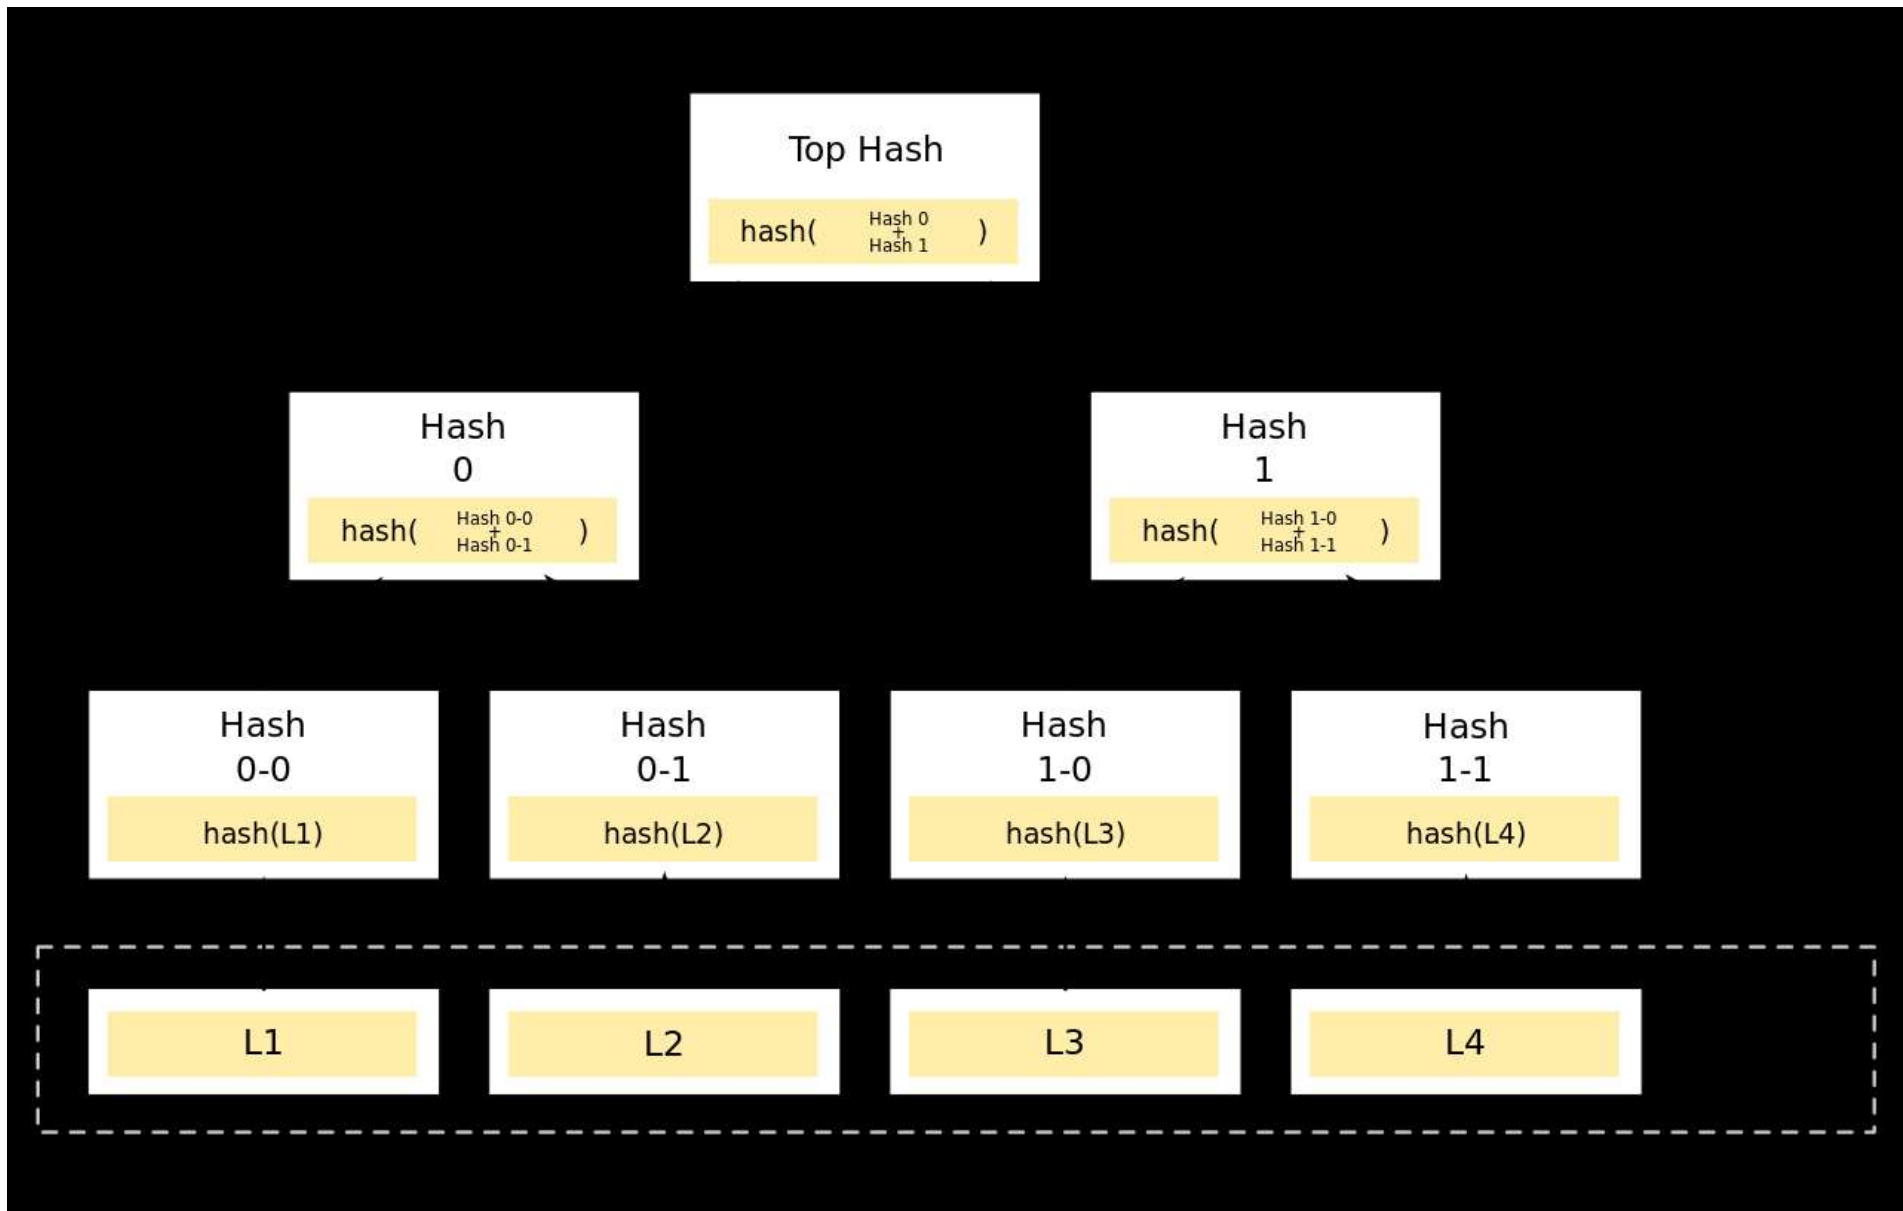
# Distributed Hash Tables

- The nodes on a P2P network can act like the buckets in a Hash Table and store records, documents
- Each P2P node generates a node-id
- Arrange the links in the P2P network so that it is easy to locate a node
- Insert content into the network by storing in the node that has a node id that is 'nearest' to the hash of the document



# Merkle Tree

- A tree in which every leaf node is labelled with the hash of a data block
  - Every non-leaf is labelled with the hash of the labels of its child nodes
  - Diagram on next slide shows a binary hash tree – but can be of arbitrary degree
- 
- Enables entire tree to be verified easily
  - Can easily prove that a block belongs to a tree
  - If you want to prove that a block is included in a tree – can provide its hash and the Merkle-Path (or hashes) that takes you to the root




By Azaghal - Own work, CC0, <https://commons.wikimedia.org/w/index.php?curid=18157888>



# Proof-of-Work

- Initially proposed in 1997 by Adam Back in his HashCash system
- Trying to replicate the concept of a 'stamp' in Email
- Construct the info contents of stamp: recipient, date, version etc
- Append a random string – Hash the resulting string
- Compute a SHA-1 hash – see if the first 20 bits(5 hex digits) are 0
- If not, increment the string and try again – on avg  $2^{20}$  tries:
- Quick to verify

X-Hashcash: 1:20:1303030600:adam@cypherspace.org::McMybZlhxKXu57jd:ckvi



# PoW is the basis for Bitcoin Mining

- Now a huge industry in China, Russia, Iceland and elsewhere
- GPUs and ASICs calculate PoW on each Bitcoin block in order to gain the 'Block Reward'
- 'Difficulty' – number of zeros on hash – is dynamically adjusted



# Public/Private Key Systems

- First disclosed (invented?) in 1976 by Diffie & Hellman
- Involves creating two linked keys – one that is made public (PK) and another that is kept secret.
- Either Key works such that
  - $\text{Key ( Plaintext Message )} = \text{Ciphertext Message}$   
[ you can encrypt the message with either key]
- Applying one key reverses the effect of the other
  - $\text{PK (Plaintext)} = \text{Ciphertext}; \text{SK (Ciphertext)} = \text{Plaintext}$
  - $\text{SK (Plaintext)} = \text{Ciphertext}; \text{PK (Ciphertext)} = \text{Plaintext}$
- If  $\text{PK}_A$  is public, Anyone can read  $\text{SK}_A(\text{Plaintext})$  but only A could have produced it (signature)
- Anyone can produce  $\text{PK}_A(\text{Plaintext})$  but only A can read it (Encryption/Enveloping)

# Two Commonly Used Systems

- Rivest, Shamir & Adleman (RSA)
  - A key pair generation process yields the 2 related keys
  - Security level determined by key length: 512, 1024, 2048... currently RSA-2048 is recommended for good security in a 30-year time-frame
- Elliptic Curve Cryptography (ECC)
  - First proposed in 1985 but did not enter wide use until 2004/5
  - Generate a private key at random – can derive the public key

# Digital Signatures & Identities

- Signing a message is done by applying a HASH function to the message – generating the message digest
- Apply a private key to the Digest to produce a signature
- Users can generate identities just by inventing key-pairs
- Sometimes it is useful to associate a “Name” with a public key
- Can be done on a pair-wise basis – Key Signing party!
- Can also be done with a certificate - Public Key Infrastructure (PKI)

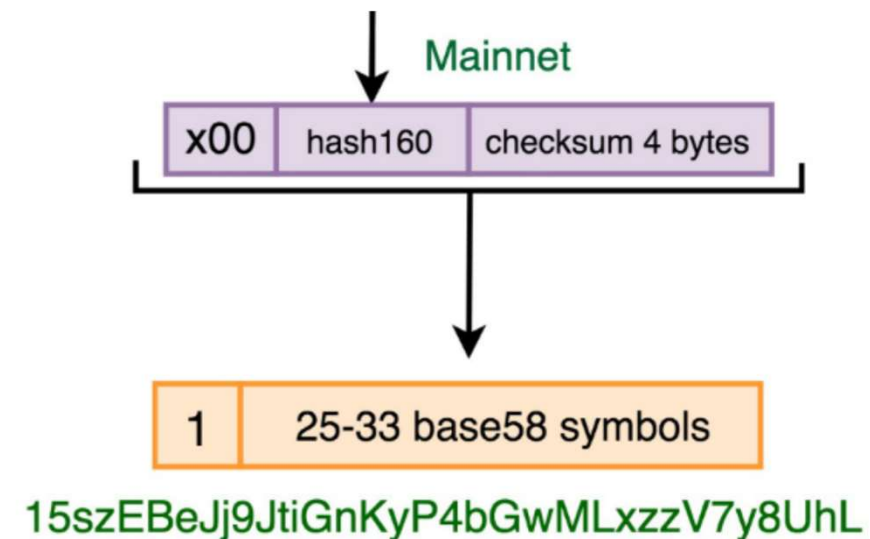
Name: Donal O'Mahony PubKey: 0x567fda3c6 DateofIssue ValidUntil:	Signature of Trusted Entity
--	-----------------------------

# Cryptographic Payment - Bitcoin

- Most electronic payments require a Trusted Intermediary (e.g. Bank, CC company, Revolut, Paypay)
- Bitcoin was described in 2008 by Satoshi Nakamoto in his paper “Bitcoin: A Peer-to-Peer Electronic Cash System”
- Set out to create an Internet-wide, decentralized, payment system that did not require any intermediaries

# Identities & Transactions

- Bitcoin users come into being when a user generates a key pair
  - The public key is hashed and the lower 160 bits is pre-pended with x00 and a 4-byte checksum added – this is the 'Bitcoin Address'
  - There is no registration
  - Initially, the user has no Bitcoin, but anyone can send him some if provided with his public key
- Once the user has acquired some (from a friend or an exchange), he can send it to other bitcoin addresses



# Transactions

- Bitcoin transactions involve 'coins' that are referred to as Unspent Transaction Outputs (UTXOs)
- Each transaction has references a list of inputs(UTXOs) and a list of outputs
- In the simplest case, Alice has a single UTXO worth 10BTC and wants to send it to Bob

Alice's 10BTX UTXO	Alice's Signature	Bob's new UTXO
--------------------	-------------------	----------------

- When this is written to the Ledger, Bob has a new 10BTC UTXO and Alice's UTXO is 'spent'



# Transactions....

- Bob can use his new 10BTC UTXO in future transactions by referencing this newly created output
- Under the hood, the signature verification is implemented with Bitcoin Script – this is a set of simple stack-based primitives that are executed by anyone who needs to test the validity of the transaction
- If Alice wants to pay Bob 7 BTC, but she only has a 10BTC UTXO, she can create a transaction with 2 outputs - 7BTC for Bob and another output that sends 3BTC back to herself – she might well create a new address for this!
- The sum of the inputs should (almost) match the sum of the outputs – a small surplus of input over output can be pocketed by the entity that verifies the transaction as a transaction fee

# Putting Transactions into Blocks - Mining

- Alice launches her transaction by sending it over the P2P network – where it will be delivered to every Bitcoin node in the network - any of the nodes can engage in mining
- Miners build a pool into which each new transaction is added
- Miners select transactions from this pool and pack them into a potential new block
- They include the block hash of the last block in this new one – so that the blocks link into a block-chain or ledger
- They add a special 'Coinbase' transaction which pays them the 'Block Reward' of (currently) 12.5BTC
- Then they engage in a Proof-of-Work exercise to find a nonce that generates a block hash with the required difficulty.
- If they find a valid block – they broadcast it to all other nodes
- All miners try to build on top of the longest chain
- The block reward halves every 210,000 blocks (4 years @ 1block/10mins) – started @50 BTC – now 6.25 - goes to zero after 21 million coins produced (in 2140)

# Observing the Bitcoin Blockchain

- Blockchain explorers like blockchain.com allow visibility into what is happening on the blockchain
- <https://www.blockchain.com/explorer>

# Bitcoin Limitations

- Bitcoin involves every miner in the world doing (almost) exactly the same computations to achieve a result (the ledger) which represents a “Distributed Consensus” on what the ledger should be – It does this with no controlling middle-man
- As a ledger (database) it is extremely slow and inefficient
- One block is produced every 10 minutes
- With a median to average transaction size – this allows 3.3 to 7 Transactions/second
- Bitcoin script is extremely limited in implementing logic beyond a simple address-to-address transfer